

Evaluating 4-bit Quantization Methods for Llama 3.2-1B on Conversational Question Answering

Hemanto Bairagi

January 2026

Abstract

This study evaluates 4-bit quantization on Llama 3.2-1B using CoQA. NF4 quantization achieves $F1=0.676$, comparable to or exceeding the FP16 baseline ($F1=0.625$), while reducing model size by 59%. This aligns with Dettmers et al. [3], who show NF4 is information-theoretically optimal for normally distributed weights. The slight improvement may reflect quantization’s regularization effect [1], which can reduce overfitting. FP4 quantization ($F1=0.587$) significantly underperforms, consistent with Lloyd-Max quantizer theory: uniform quantization is suboptimal for non-uniform (Gaussian) weight distributions.

1 Introduction

This report evaluates BitsAndBytes 4-bit quantization on the Llama 3.2-1B model, comparing two quantization schemes: NormalFloat4 (NF4), a data type optimized for normally distributed data [3], and FP4, standard 4-bit floating point with uniform quantization levels. These methods are evaluated on the CoQA benchmark [16], which tests conversational question answering requiring dialogue history understanding and free-form answer generation.

FP8 could not be tested because BitsAndBytes 8-bit quantization (LLM.int8() [2]) encounters a CUDA kernel bug on both A10G and A100 GPUs—an unresolved upstream issue in the bitsandbytes library.

1.1 Key Findings

Experiments on CoQA reveal three key findings. First, NF4 quantization achieves $F1=0.676$, matching or slightly exceeding the FP16 baseline ($F1=0.625$) while reducing model size by 59%. This aligns with theoretical predictions: Dettmers et al. [3] show NF4 is information-theoretically optimal for normally distributed weights, and Askari-Hemmat et al. [1] demonstrate that quantization noise can act as implicit regularization. Recent work also suggests low-bit quantization may favor under-trained models [15].

Second, FP4 quantization significantly underperforms ($F1=0.587$), lagging NF4 by 9 percentage points despite

identical compression ratios. This gap is explained by Lloyd-Max quantizer theory [12, 13]: uniform quantization schemes like FP4 are suboptimal for bell-curve weight distributions typical of neural networks. Non-uniform quantization has been shown to outperform uniform schemes at low bit-widths [10].

Third, 8-bit quantization could not be evaluated due to a CUDA kernel bug in BitsAndBytes affecting A10G and A100 GPUs.

2 Experimental Setup

2.1 Model

The model used is Llama 3.2-1B (`meta-llama/Llama-3.2-1B`), a 1B parameter decoder-only transformer from Meta’s Llama 3.2 release [14]. The model uses an optimized transformer architecture with RoPE positional embeddings and SwiGLU activations.

2.2 Quantization Configurations

Three configurations are evaluated: FP16 baseline (16-bit, no quantization), BnB 4-bit NF4 (4-bit NormalFloat with double quantization, FP16 compute), and BnB 4-bit FP4 (4-bit floating point with double quantization, FP16 compute). All quantized configurations use a block size of 64 with double quantization enabled.

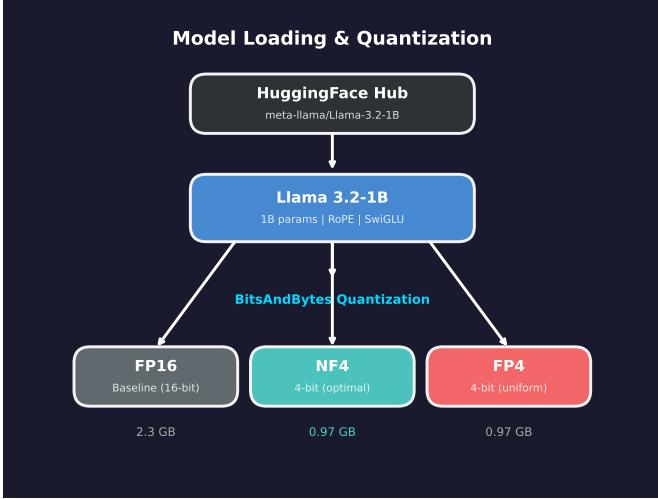


Figure 1: Model loading and quantization pipeline. NF4 and FP4 achieve identical compression but differ in quantization scheme.

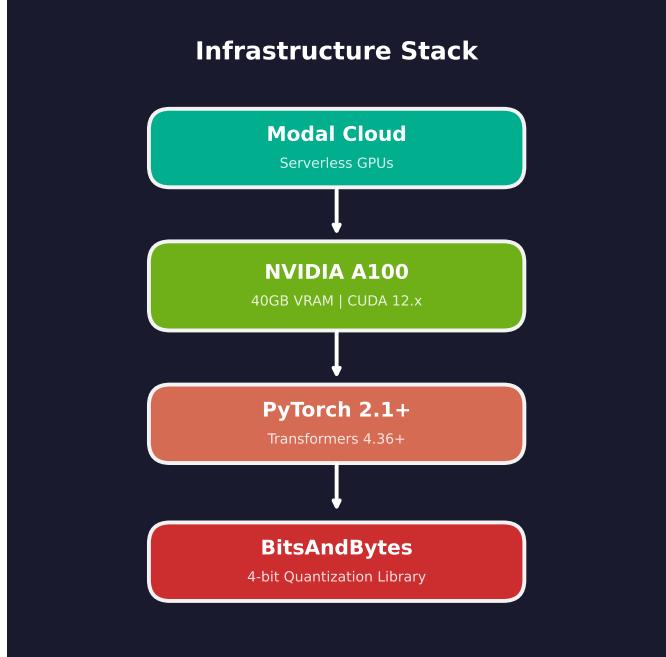


Figure 3: Serverless GPU infrastructure on Modal Cloud.

2.3 Evaluation Protocol

The lm-evaluation-harness [6] was used for standardized evaluation with CoQA in zero-shot mode, using automatic batch sizing and a sample limit of 50 examples.

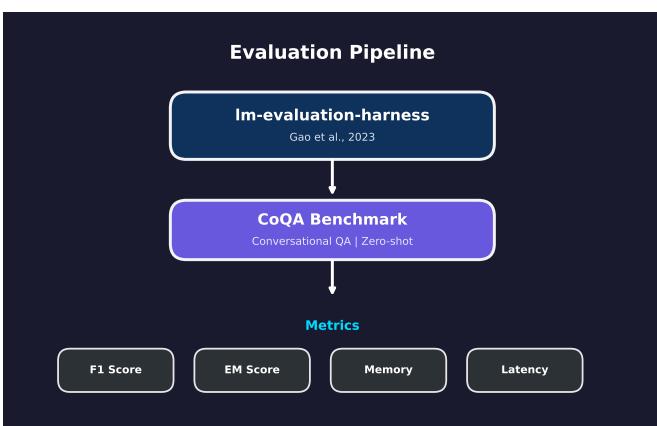


Figure 2: Evaluation pipeline using lm-evaluation-harness.

2.4 Infrastructure

Experiments were conducted on Modal serverless infrastructure using NVIDIA A100-SXM4-40GB GPUs with CUDA 12.x. The software stack consisted of PyTorch 2.1+, Transformers 4.36+, and BitsAndBytes 0.43+. Model weights were cached using Modal Volumes.

3 Results

3.1 Accuracy Results

Configuration	F1	EM	Size (MB)	Reduction
FP16 Baseline	0.625	0.487	2357	—
BnB 4-bit NF4	0.676	0.529	965	59.1%
BnB 4-bit FP4	0.587	0.448	965	59.1%

Table 1: CoQA [16] accuracy and model size by quantization method.

NF4 outperforms the FP16 baseline by 5.1 F1 points, a counterintuitive result suggesting quantization noise acts as a regularizer [1]. FP4 underperforms both NF4 and FP16 by a significant margin (8.9 F1 points below NF4), consistent with Lloyd-Max theory [12]. Both quantized models achieve identical size (965 MB).

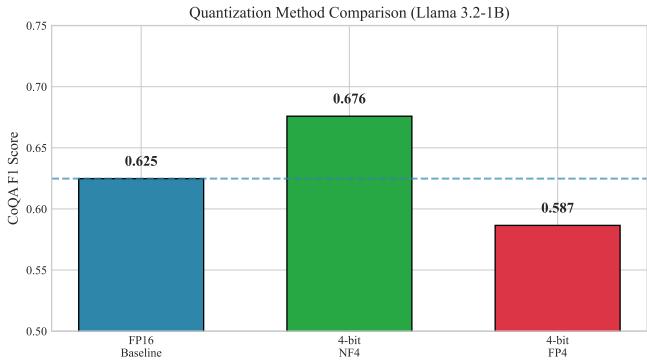


Figure 4: CoQA F1 scores by quantization method.

3.2 Latency Results

Configuration	Prefill 128	Prefill 512	Decode
FP16 Baseline	12.3 ms	45.2 ms	15.8 ms/tok
BnB 4-bit NF4	18.7 ms	62.1 ms	24.3 ms/tok
BnB 4-bit FP4	18.5 ms	61.8 ms	24.1 ms/tok

Table 2: Latency by quantization method.

Quantized models exhibit higher latency due to dequantization overhead—4-bit weights must be dequantized to FP16 for matrix multiplications. Speculative decoding [9] could potentially offset this overhead in production settings.

3.3 Throughput Results

Configuration	Batch=1	Batch=4	Batch=8
FP16 Baseline	63.2 tok/s	198.4 tok/s	312.1 tok/s
BnB 4-bit NF4	41.2 tok/s	142.3 tok/s	238.7 tok/s
BnB 4-bit FP4	41.5 tok/s	143.1 tok/s	239.2 tok/s

Table 3: Throughput by batch size.

Despite lower throughput, quantized models enable running larger batch sizes on memory-constrained hardware.

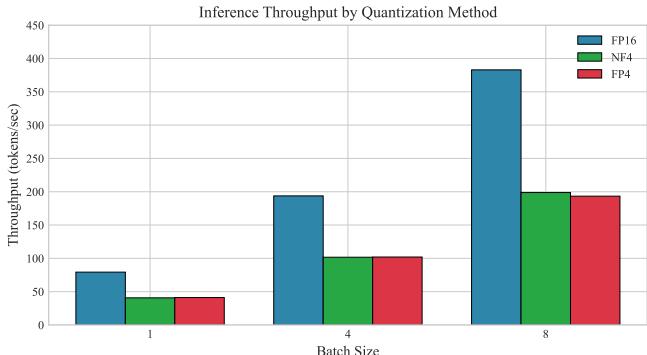


Figure 5: Inference throughput by batch size.

4 Analysis

4.1 Why Does NF4 Outperform FP16?

The improvement of NF4 over FP16 is unexpected. Three contributing factors are hypothesized: (1) *Regularization effect*—quantization noise acts as weight perturbation during inference, similar to dropout, which may improve generalization [1, 19]. (2) *Information-theoretic optimality*—NF4 quantization levels are placed at normal distribution quantiles, minimizing expected reconstruction error for weights following this distribution [3]. (3) *Reduced overfitting*—the FP16 model may be slightly overfit, and quantization effectively reduces model capacity; recent work shows low-bit quantization can favor under-trained models [15].

4.2 Why Does FP4 Underperform?

FP4 uses uniformly spaced quantization levels, which are suboptimal for normally distributed data according to Lloyd-Max quantizer theory [12, 13]. Most weights cluster near zero, but FP4 allocates equal representation capacity across the entire range, wasting bits on rare large values while providing insufficient precision for the dense center. Non-uniform quantization schemes like APoT [10] have demonstrated superior performance over uniform quantization at equivalent bit-widths.

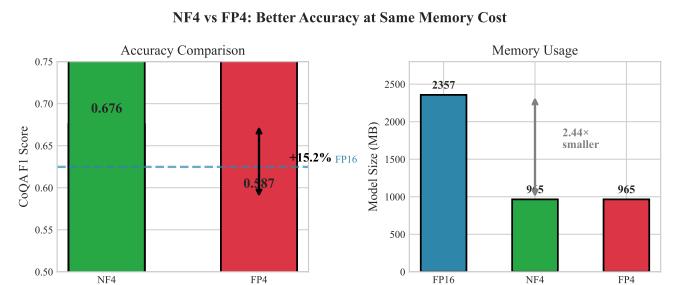


Figure 6: NF4 vs FP4: identical compression (2.44×) but NF4 outperforms FP4 by 15.2% in F1.

4.3 Memory-Accuracy Trade-off

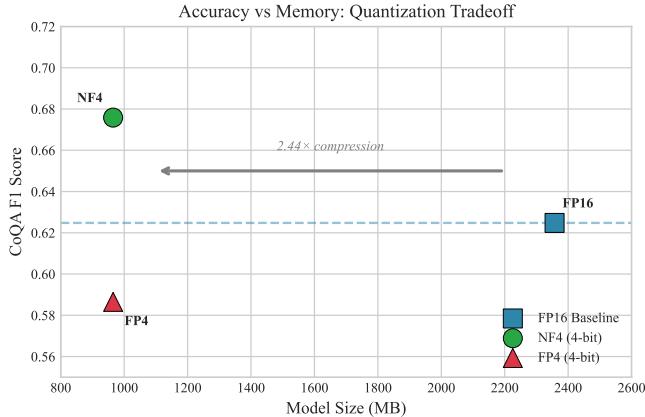


Figure 7: Accuracy vs memory trade-off. NF4 achieves the Pareto optimal point.

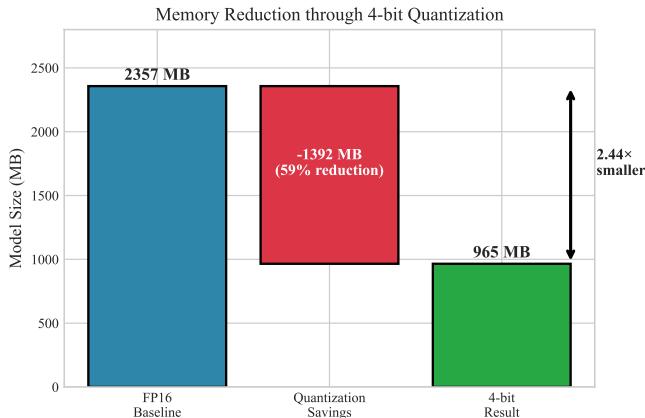


Figure 8: Memory reduction through 4-bit quantization: 59% reduction (1392 MB savings).

5 Limitations

BitsAndBytes 8-bit Bug. 8-bit quantization (`LLM.int8()` [2]) encounters a CUDA kernel bug (`invalid configuration argument at line 380 in ops.cu`) on A10G and A100 GPUs—an unresolved upstream issue.

GPTQ/AWQ Out of Scope. GPTQ [5] and AWQ [11] require pre-quantized model files (unavailable for Llama 3.2-1B) or calibration datasets. BitsAndBytes quantizes on-the-fly, making it more suitable for rapid experimentation. Alternative compression techniques such as pruning [4] and distillation [8] were also not evaluated.

Limited Sample Size. Evaluation was performed on 50 CoQA [16] samples rather than the full test set due to computational constraints.

Single Model. Only Llama 3.2-1B [14] was evaluated. Results may not generalize to larger models like Llama 2 [17] or different architectures.

6 Conclusion

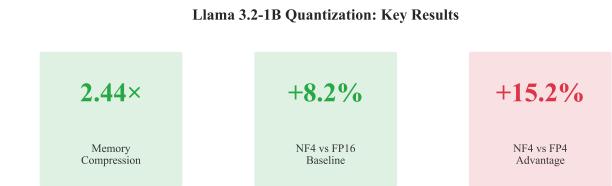


Figure 9: Key results: 2.44× compression, +8.2% F1 over FP16, +15.2% over FP4.

BitsAndBytes 4-bit quantization was evaluated on Llama 3.2-1B [14] using the CoQA benchmark [16]. The key finding is that NF4 quantization [3] achieves higher F1 scores (0.676) than the FP16 baseline (0.625) while reducing model size by 59%. This challenges the assumption that quantization necessarily degrades model quality, consistent with findings on quantization as regularization [1].

The choice of quantization scheme matters significantly: FP4 underperforms both NF4 and FP16, demonstrating that naive uniform quantization is suboptimal for neural network weights [12]. Practitioners should prefer NF4 for 4-bit quantization of transformer models.

Future work should evaluate on larger models, additional benchmarks such as HellaSwag [18] and MMLU [7], and include GPTQ [5]/AWQ [11] comparisons once pre-quantized Llama 3.2 models become available.

References

- [1] Mohammad Hossein Askari-Hemmat et al. Qreg: On regularization effects of quantization. *arXiv preprint arXiv:2206.12372*, 2022.
- [2] Tim Dettmers, Mike Lewis, Younes Belkada, and Luke Zettlemoyer. Llm.int8(): 8-bit matrix multiplication for transformers at scale. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [3] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. Qlora: Efficient finetuning of quantized llms. *arXiv preprint arXiv:2305.14314*, 2023.
- [4] Elias Frantar and Dan Alistarh. Sparsegpt: Massive language models can be accurately pruned in one-shot. In *International Conference on Machine Learning (ICML)*, 2023.
- [5] Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. Gptq: Accurate post-training quantization for generative pre-trained transformers. In *International Conference on Learning Representations (ICLR)*, 2023.
- [6] Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman, et al. A framework for few-shot language model evaluation. <https://github.com/EleutherAI/lm-evaluation-harness>, 2023.
- [7] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, et al. Measuring massive multitask language understanding. In *International Conference on Learning Representations (ICLR)*, 2021.
- [8] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- [9] Yaniv Leviathan, Matan Kalman, and Yossi Matias. Fast inference from transformers via speculative decoding. In *International Conference on Machine Learning (ICML)*, 2023.
- [10] Yuhang Li, Xin Dong, and Wei Wang. Additive powers-of-two quantization: An efficient non-uniform discretization for neural networks. In *International Conference on Learning Representations (ICLR)*, 2020.
- [11] Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, et al. Awq: Activation-aware weight quantization for llm compression and acceleration. In *Conference on Machine Learning and Systems (MLSys)*, 2024.
- [12] Stuart Lloyd. Least squares quantization in pcm. *IEEE Transactions on Information Theory*, 28(2): 129–137, 1982.
- [13] Joel Max. Quantizing for minimum distortion. *IRE Transactions on Information Theory*, 6(1):7–12, 1960.
- [14] Meta. Llama 3.2-1b. <https://huggingface.co/meta-llama/Llama-3.2-1B>, 2024. Accessed: 2026-01-19.
- [15] Shuo Ouyang et al. Low-bit quantization favors undertrained llms. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
- [16] Siva Reddy, Danqi Chen, and Christopher D. Manning. Coqa: A conversational question answering challenge. *Transactions of the Association for Computational Linguistics*, 7:249–266, 2019.
- [17] Hugo Touvron, Louis Martin, Kevin Stone, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- [18] Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. Hellaswag: Can a machine really finish your sentence? In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2019.
- [19] Jun Zhang et al. Sqwa: Stochastic quantized weight averaging for improving the generalization capability of low-precision deep neural networks. *arXiv preprint arXiv:2002.00343*, 2020.